

Author

Kaushik Harsha

kaushikharsha2020@gmql.com

I am a Data Science Enthusiast, currently pursuing B.Tech in Computer Science & Engineering. I also enjoy full stack web development.

AI/LLMs Usage

AI/LLM tools were used ~20% to assist with styling, backend jobs, caching, etc.

Description

HealthLink, the Hospital Management System is a Flask-based full-stack application designed to streamline hospital workflows such as patient registration, appointment handling, doctor availability, and treatment record management. The system supports three user roles: Admin, Doctor, and Patient, each with distinct permissions and dashboards. The backend is powered by Flask, SQLite (via SQLAlchemy ORM), Redis caching, and Celery-based asynchronous task processing for reminders, monthly reports, and CSV exports. The frontend is developed using Vue.js and Bootstrap to ensure a responsive and interactive user interface.

Technologies Used

- Python – Main Backend handler using flask and other backend jobs
- Flask – Backend Server
- Flask-SQLAlchemy – Backend Database Management
- Flask-Security – Authentication and RBAC
- SQLite3 - Database
- Redis – For Caching and running celery tasks
- Celery – Running asynchronous tasks using beats and workers on redis server
- Vue.js – Main Frontend Framework
- Bootstrap – For styling and responsive design(PWA)
- HTML – For templating monthly PDF reports for doctors
- JavaScript – For client side services
- WSL2 – For running the Redis server and Celery workers(Backend)

DB Schema Design

The database schema consists of the following tables:

- users (id, name, email, password, fs_uniquifier, active)
- roles (id, name, description)
- user_roles (user_id, role_id)
- specializations (id, name, description)
- doctors (id, u_id, specialization_id, availability, contact_number)
- patients (id, u_id, age, gender, contact_number, address, emergency_contact)
- appointments (id, doctor_id, patient_id, appointment_date, appointment_time, status)
- treatments (appointment_id, diagnosis, tests_done, prescription, notes)

Design Rationale:

The schema follows a normalized relational model ensuring clear relationships between users, doctors, patients, appointments, and treatments. Cascade rules are used for maintaining integrity across dependent entities.

API Design:

- GET - /users
- GET - /users/<id>
- GET - /doctors
- GET - /doctor?id=<id>
- GET - /doctor?uid=<uid>
- GET - /patients
- GET - /patient?id=<id>
- GET - /patient?uid=<uid>
- GET - /appointments
- GET - /specialization?id=<id>
- GET - /specialization?name=<name>

- POST - /auth/login
- POST - /auth/register
- POST - /doctors
- POST - /appointments
- POST - /treatments
- POST - /specializations

- PUT - /users/<id>

- PATCH /users/<id>
- PATCH /doctors/<id>
- PATCH /patients/<id>
- PATCH /appointments/<id>

- DELETE /users/<id>
- DELETE /patients/<id>
- DELETE /appointments/<id>

Features Implemented

- Role-based login for Admin, Doctor, and Patient.
- Admin dashboard displaying total doctors, patients, and appointments.
- Doctor dashboard showing upcoming appointments and patient lists.
- Patient dashboard showing available doctors, booking interface, and history.
- Doctor availability management for 7 days.
- Appointment booking, cancellation, and rescheduling logic.
- Treatment recording with diagnosis, notes, and prescriptions.
- Redis caching of common queries for speed optimization.
- Celery-based scheduled reminders and monthly reports.
- CSV export background job for patient treatment history.

Video

https://drive.google.com/file/d/1Ri6pbTvTVafEx6gPJZcFDz353iEkWJca/view?usp=drive_link